

*Exceptional service in the national interest*



OGC | Open Grid Computing, Austin, TX



# LDMS Version 4.3 Tutorial Part 1: Basics

<https://github.com/ovis-hpc/ovis>

Jim Brandt, Ann Gentile, Ben Allan

Sandia National Laboratories

Tom Tucker

Open Grid Computing, Inc.

SAND2017-5153 O



Sandia National Laboratories is a multi-mission laboratory managed and operated by National Technology and Engineering Solutions of Sandia, LLC., a wholly owned subsidiary of Honeywell International, Inc., for the U.S. Department of Energy's National Nuclear Security Administration under contract DE-NA-0003525.

04/2017

# Advance Set-up (site specific: wifi UCF\_Guest)

We will be using virtual machines hosted at Open Grid Computing

```
$ ssh user<#>@ldmscon.ogc.us
```

```
$ password: user<#>
```

```
$ ssh user#@compute<#>
```

```
$ password: user<#>
```

# Schedule

- **9:00 – 10:30 Instruction (Basics)**
- **10:30 – 10:45 Break**
- **10:45 – 12:30 Instruction (Basics)**
- 12:30 – 1:30 Lunch
- 1:30 – 3:00 Instruction
- 3:00 – 3:15 Break
- 3:15 – 5:00 Instruction

# Tutorial Format (Basic)

## Overview of the Lightweight Distributed Metric Service (LDMS) (9 slides)

- Overview of the LDMS framework
- LDMS architecture description

## Setup (3 slides)

- Environment setup description and verification
- Introduction to support programs and helper scripts for use in lab work

## Hands-on exercises, instructor walk through, and facilitated student exploration:

- **Exercise 1:** Memeater (1 slide)

### *Configuring and deploying a distributed monitoring system with storage*


- **Exercise 2:** Configuring and Running Samplers (37 slides – ~1 hour)
  - Sampler startup and local and remote verification
  - Intro to ldmsd\_controller and ldms\_ls
- **Exercise 3:** Configure Aggregators (13 slides – ~30 min)
  - Aggregation startup and verification using local samplers
  - Aggregation of all other attendees' (remote) samplers
- **Exercise 4:** Aggregating From Remote Hosts: Building a Distributed Monitoring System (4 slides – ~45 min)
- **Exercise 5:** Dynamic Configurations and Resilience (4 slides – ~20 min)
- **Exercise 6:** Storing Data In CSV Format (8 slides – ~20 min)

# LDMS Overview

- What is the Lightweight Distributed Metric Service (LDMS)?
  - Daemon based data sampling
    - Collect numeric data
  - Move and aggregate data
  - Store data
  - Analyze data
    - Troubleshooting
    - Optimization
    - Inform future designs
- Typical use cases
  - Identify applications memory (and other resource) utilization behaviors
  - Identify network congestion
  - Determine over-provisioned resources
  - Identify heavy Lustre users

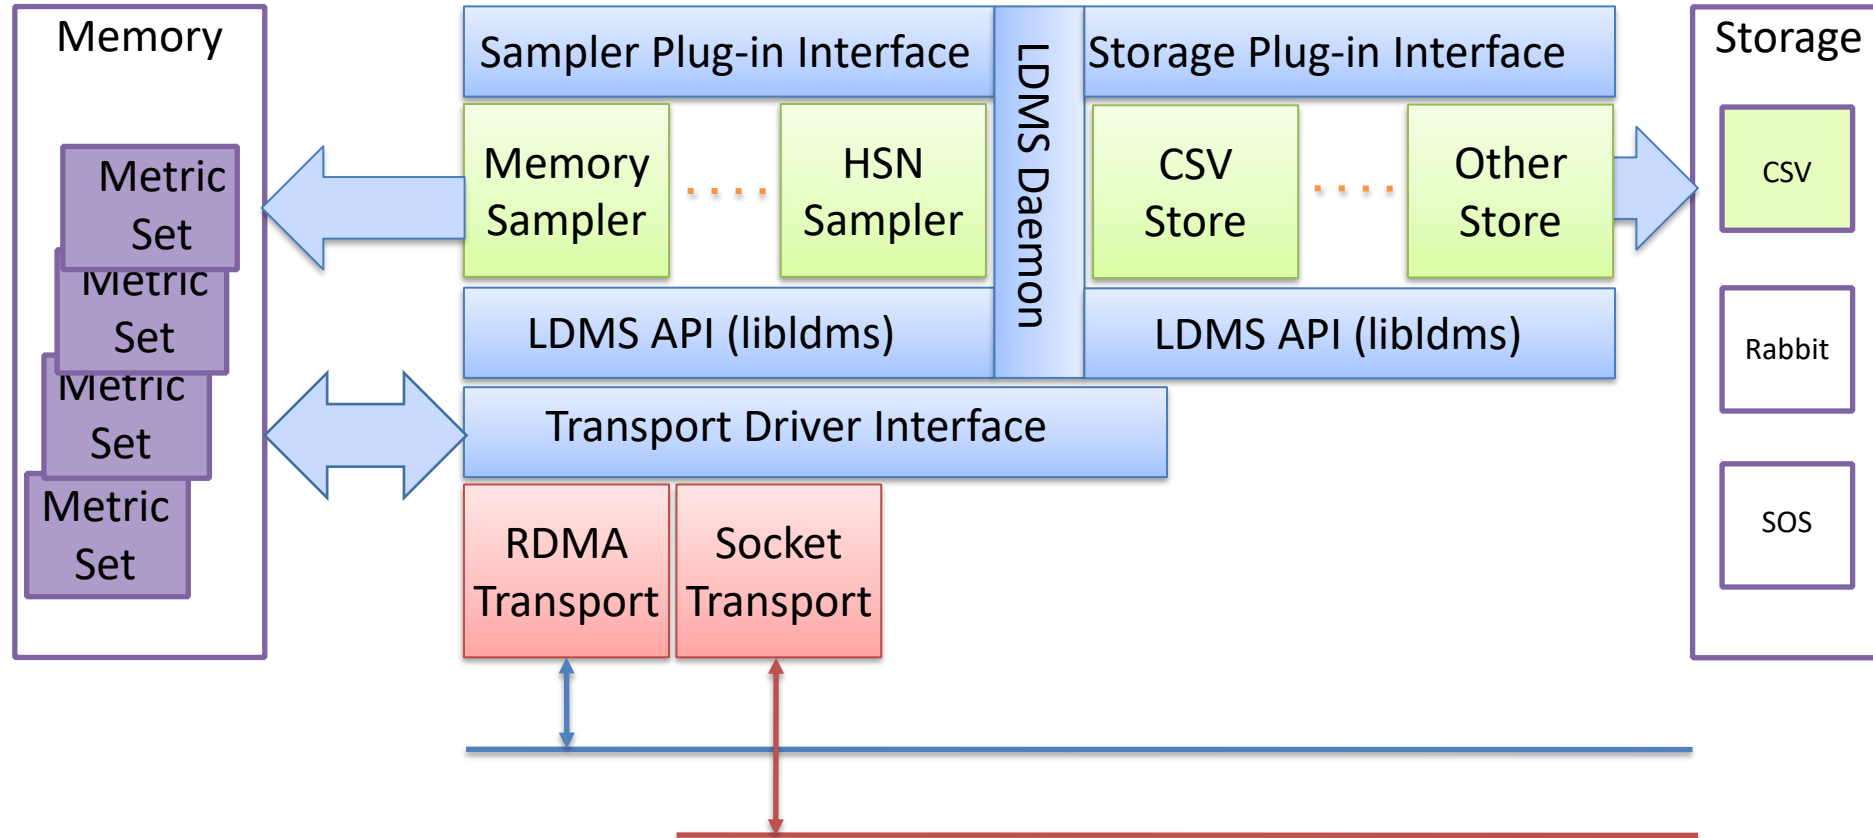
Feature	Generic Monitoring Systems (e.g., Nagios, Ganglia)	Vendor Specific Monitoring Systems (e.g., Cray SEDC)	Application Profilers (e.g., CrayPat, OpenSpeedShop)	LDMS
Scalability	Low	Medium	Low	High
Typical sampling frequency	Low	Low	High	Medium
Overhead	Medium	Low	High	Low
Ease of adding new base collectors	High	Low	Medium	High
Ease of using data in generic analysis tools	High	Low	Low	High
Suitability of data for researchers (e.g., system performance analysis, resilience)	Low	Low	Low	High
Suitability of data for use in dynamic feedback (e.g., application or system software can access run-time data for reconfiguration)	Low	Medium	Low	High
Portable	Yes	No	Yes	Yes
Ability to provide coherent system snapshots	No	No	No	Yes
Ability to provide information obtained under production conditions	Yes	Yes	No	Yes
Can provide run-time information to platform users	Yes	Yes	No	Yes

# Lightweight Distributed Metric Service (LDMS) High Level Overview



\* Only the current data is retained on-node

# LDMS Plugin Architecture





# Metric Set Memory

## Metric Meta Data

- Generation Number

### Metric Descriptor

- Name
- Component ID
- Type
- Offset

### Metric Descriptor

- Name
- Component ID
- Type
- Offset

### Metric Descriptor

- Name
- Component ID
- Type
- Offset



## Metric Data

- Meta Data Generation Number
- Data Generation Number
- Consistent Status

Value

Value

Value



# Resources

- Documentation (Building, Using)
  - <https://github.com/ovis-hpc/ovis/wiki>
- Source Code
  - <https://github.com/ovis-hpc/ovis>
  - git clone <https://github.com/ovis-hpc/ovis.git>
  - git branch -a # Will show all available branches
  - git branch -a | grep "\-4.3" # Will show all version 4.3 branches
  - git checkout -b OVIS-4.3.<x> origin/OVIS-4.3.<x> # Will check out branch origin/OVIS-4.3.<x> under the name OVIS-4.3.<x>
  - git branch # Will show currently checked out branch
- Publications:
  - <https://ovis.ca.sandia.gov>
- How you can contribute
  - Post an issue at: <https://github.com/ovis-hpc/ovis/issues>
- Support
  - Bug reporting and questions: Post an issue at: <https://github.com/ovis-hpc/ovis/issues>
  - Development services: contact [tom@ogc.us](mailto:tom@ogc.us)
  - Support services: contact [tom@ogc.us](mailto:tom@ogc.us), [ldms@sandia.gov](mailto:ldms@sandia.gov)

# Supported platforms and networks

- Linux support
  - RHEL 6 and 7
  - SLES 11 & 12
  - Ubuntu
- Vendor hardware platforms running supported software
  - Cray XE6, XK and XC
  - Generic Linux clusters
  - IBM P8 & P9 (both big and little endian)
- Transports
  - Socket
  - Cray ugni
    - Aries
    - Gemini
  - RDMA
    - Infiniband
    - iWarp
    - libfabric

# Build dependencies

- Typical compute node environment
  - Autoconf  $\geq 2.63$ , automake, libtool (collectively called autotools)
  - OpenSSH-devel
  - libpapi-devel for papi and syspapi samplers
  - libpfm-devel for syspapi sampler
  - libfabric-devel if applicable transport available
- End use hosts (monitor cluster, special aggregation hosts, etc.)
  - Python 3.x
  - Swig 2.0.x
  - Doxygen for documentation
  - Cython needed for SOS
    - Get from pip
  - libcurl & libcurl-devel if using influx\_store

# LDMS Installation methods

- Manually build and install using autoconf and automake
- Deployment using RPMs

**Note1:** For this tutorial, LDMS is pre-installed on student VMs in /opt/ovis

**Note2:** We will be building and installing to local directories and will use the pre-installed software for all other exercises

# Setup

# Getting started: Log in and set up your environment

```
ssh user<#>@ldmscon.ogc.us
```

```
user<#>@ldmscon.ogc.us password: user<#>
```

```
$ ssh user<#>@compute<#>
```

Note: “/home/<user>/exercises/ldms/env/ldms-env.sh” is used to set up LDMS environment. You may need to create this file first.

You will want at least 2 terminal windows up for the tutorial

# VM directory structure

- VMs include source code, scripts and configuration files for every exercise, helper mini-applications for use in the exercises, and supporting visualization tools (e.g., gnuplot).
- Directory structure:

<code>/home/&lt;user&gt;/exercises/</code>	<code># Location of exercise related directories</code>
<code>/home/&lt;user&gt;/exercises/ldms/conf/</code>	<code># Exercise configuration files</code>
<code>/home/&lt;user&gt;/exercises/ldms/ data/</code>	<code># LDMS data</code>
<code>/home/&lt;user&gt;/exercises/ldms/ env/</code>	<code># Scripts to configure environment variables</code>
<code>/home/&lt;user&gt;/exercises/ldms/ scripts/</code>	<code># Helper scripts for deploying LDMS daemons</code>
<code>/home/&lt;user&gt;/exercises/ldms/ code/</code>	<code># memeater code</code>
<code>/home/&lt;user&gt;/exercises/ldms/ logs/</code>	<code># Place to write log files</code>
<code>/home/&lt;user&gt;/exercises/ldms/ run/</code>	<code># symlink to /tmp/run – place to write pid files</code>



# Getting started: Set up and verify your environment

- Edit environment configuration file (ldms\_env.conf)

```
OVIS_HOME=/opt/ovis

#System environment variables
export PATH=${OVIS_HOME}/bin/:${OVIS_HOME}/sbin/:${PATH}
export LD_LIBRARY_PATH=${OVIS_HOME}/lib/:${LD_LIBRARY_PATH}
export PYTHONPATH=${OVIS_HOME}/lib/python3.6/site-packages/:${PYTHONPATH}

#LDMS environment variables
export ZAP_LIBPATH=${OVIS_HOME}/lib64/ovis-ldms
export LDMSD_PLUGIN_LIBPATH=${OVIS_HOME}/lib/ovis-ldms
```

- Source your environment configuration file  
\$ source ldms\_env.conf

\*A live example of these commands can be found here:  
[Verify Environment Variables](#)

# Exercise 1: Memeater

# Compile Test Code: memeater.c

- Memeater code repeatedly allocs memory.
- Run in conjunction with LDMS to see changes in memory utilization values reported in /proc/meminfo.
- Code is located at: /home/<user>/exercises/ldms/code/**memeater.c**

```
$ cd /home/<user>/exercises/ldms/code/memeater.c
```

Compile with cc: `cc -o memeater memeater.c`

Sleep between alloc. Change  
this wrt sampling frequency.

Periodically increase  
memory allocated

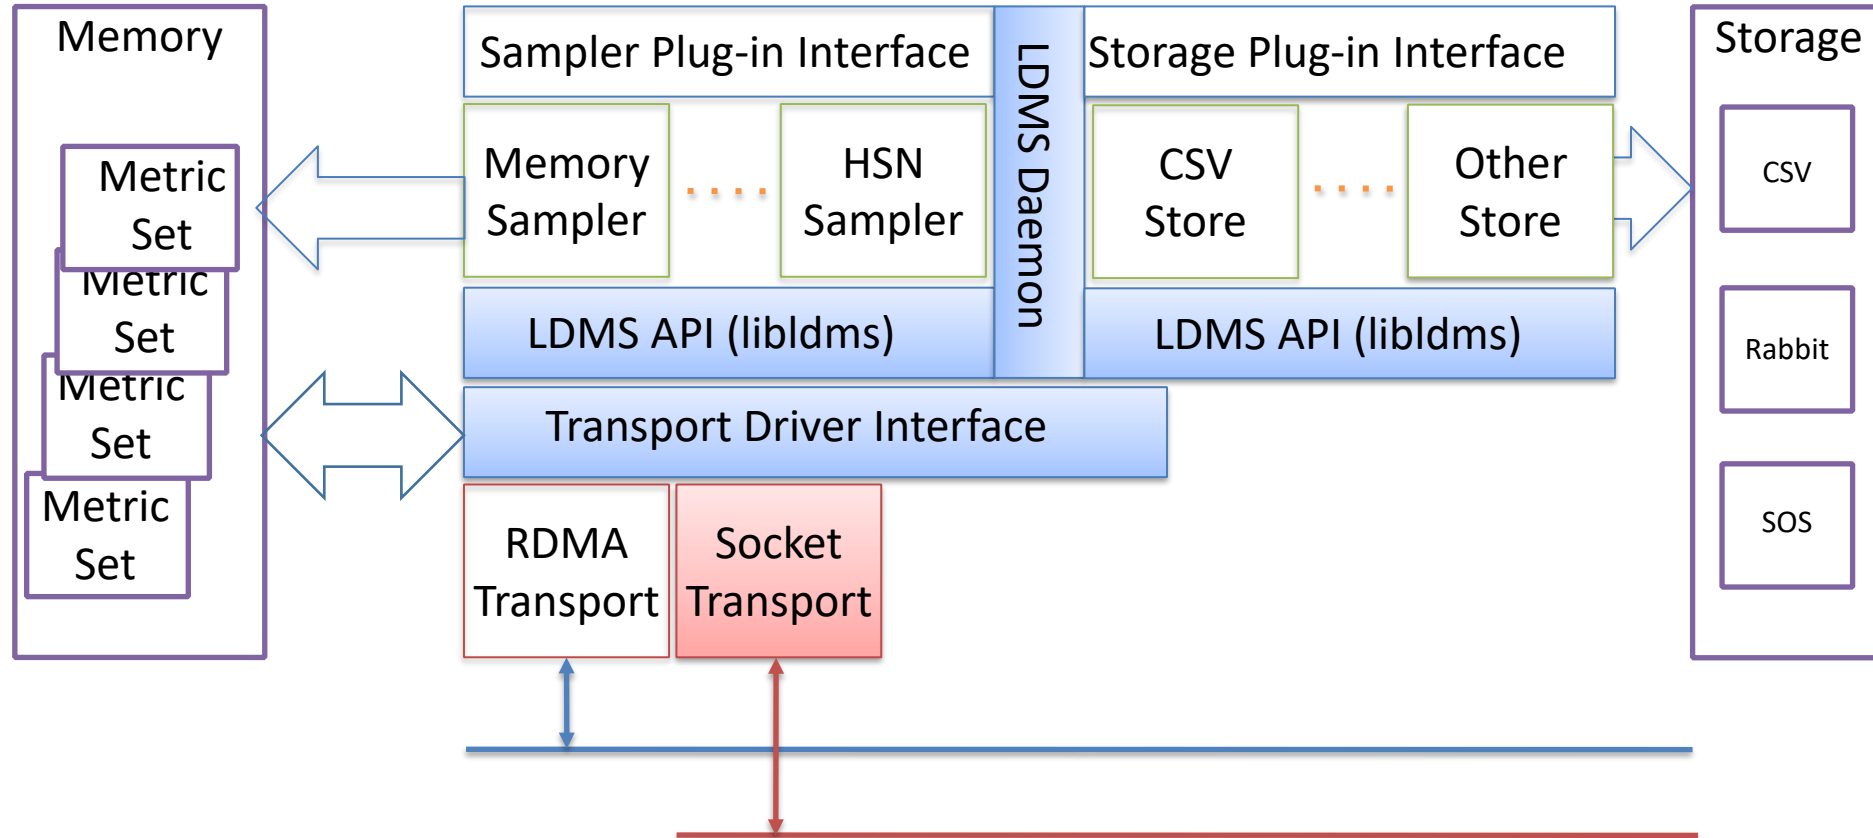
```
while (1){  
    sleep(2);  
  
    temp = (int*) realloc (keep, ((6144*6144)+count)*sizeof(int));  
    if (!temp){  
        printf( "Cannot realloc\n");  
        break;  
        /* malloc will return NULL sooner or later, due to lack of memory */  
    }  
    ...  
}  
printf("sleeping before exiting\n");  
sleep(60);  
free(keep);  
return 0;
```

Sleep before releasing memory

```
./memeater  
Active:          231148 kB  
alloc: 37748736  
  
adding 1944999541  
Active:          378616 kB  
alloc: 75497472  
...  
adding 347488691  
Active:          1263360 kB  
alloc: 301989888  
  
adding 1514442648  
adding 1528811800  
adding 1877058034  
Problems with pipe: Cannot allocate memory  
sleeping before exiting
```

# Exercise 2: Configuring and Running Samplers

# LDMS Plugin Architecture



# Start and Configure a LDMS Daemon

## Exercise Goals:

- Basic LDMS daemon startup and configuration flags/args
  - Manual and run-time configuration options
  - Output options
    - Log files and
  - man pages
    - `man ldmsd` – displays `ldmsd` man pages
    - `man ldmsd_controller` – displays “`ldmsd_controller`” man pages
- Use of `ldms_ls` utility as a diagnostic tool
  - man pages
    - `man ldms_ls` – displays `ldms_ls` man pages

# Start a LDMS daemon

- Start ldmsd with minimum configuration

```
ldmsd -x sock:10001 -l /home/<user>/exercises/ldms/logs/sampler1.log
```

- **-x**: Transport : listening port
- **-l**: Specify the log file path and name(this is not strictly necessary)

## NOTES:

- If you receive a “permission denied” error in the “sampler1.log” file, you will need to add “**-r** ldmsd.pid” at the end of the ldmsd command.
  - **-r**: The path to the pid file. Please review man page “/ldms/man/ldmsd.man” for more information
- Commands should be **written** in the command prompt window. Copy and paste may cause unnecessary issues with the command line interface

# Check Idmsd Running Status

- Using ps

```
ps auxw | grep ldmsd | grep -v grep
```

- Returns something like:  
"ovis\_pu+ 3582 0.0 0.1 401604 2204 ? ss1 12:51 0:00  
ldmsd -x sock:10001" **if running**
- Returns: **blank line if not running**

- Using ldms\_ls

```
ldms_ls -h localhost -x sock -p 10001
```

- Returns: "Connection failed/rejected." if ldmsd specified does not exist or authentication fails
- Returns: blank line if the ldmsd specified exists but has no metric sets configured
- **Remote: Replace "localhost" in the command above with another student vm e.g., "ovis-demo-28"**

- Also check ports

```
netstat -an | grep 000
```

```
tcp        0      0 0.0.0.0:10001 0.0.0.0:*    LISTEN
```

- Troubleshooting: Also check the log for clues if operation seems wrong.



## EXAMPLE: Start and Check LDMS Daemon

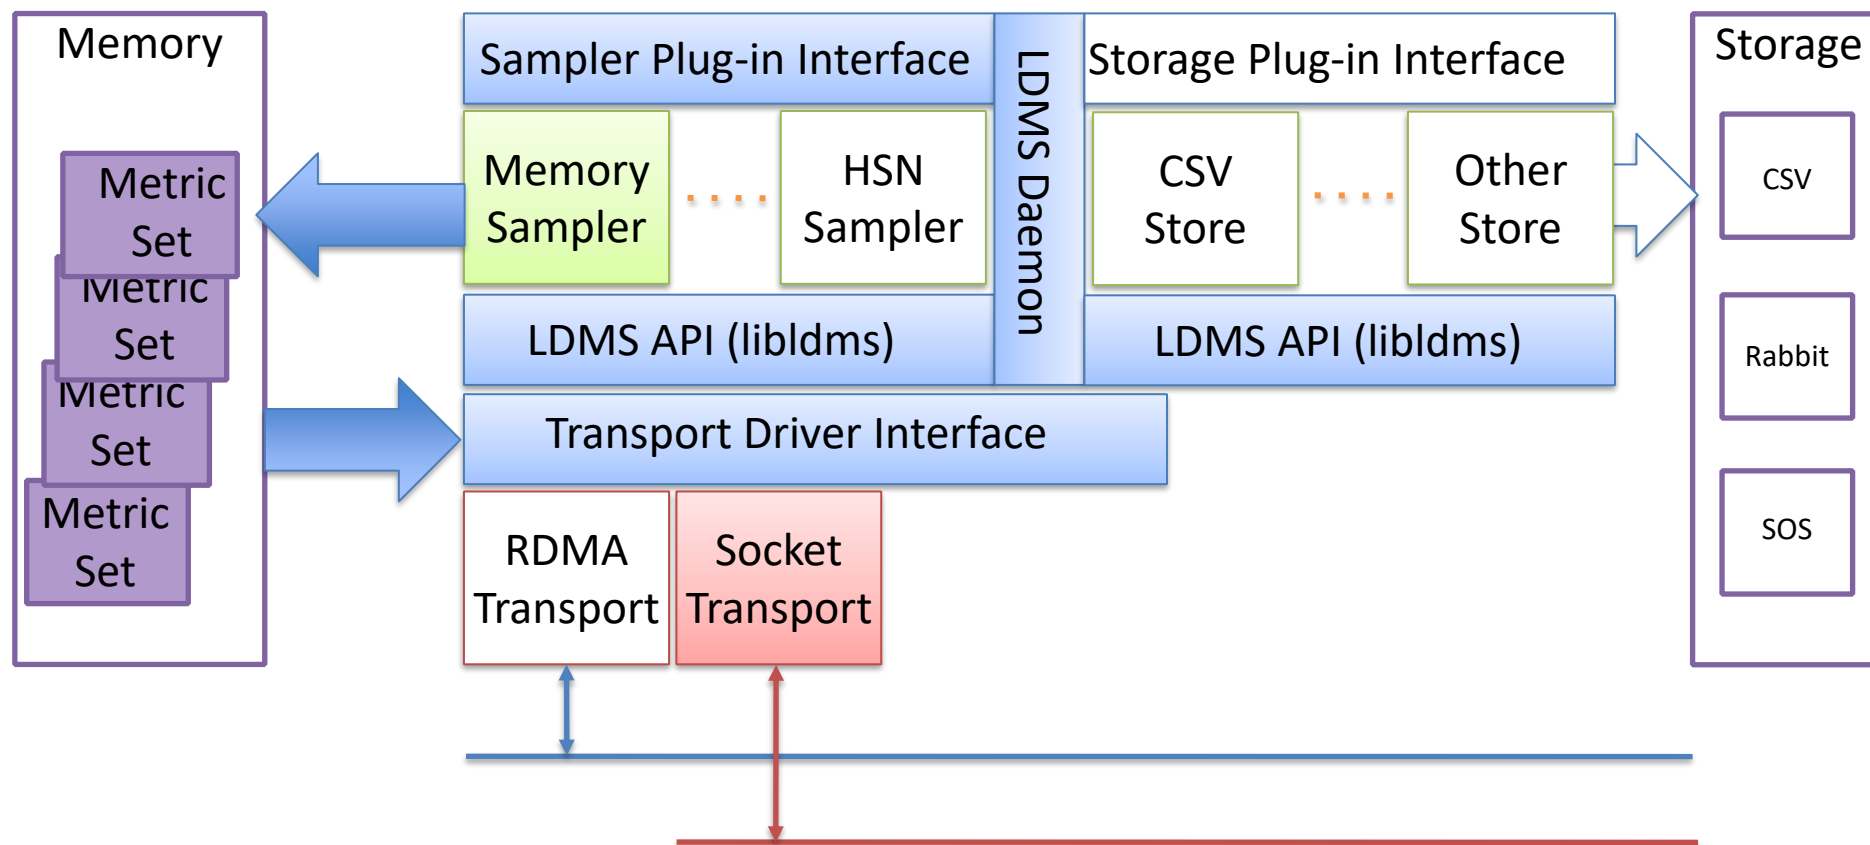
Please see the [Start and Check an LDMS daemon](#) to view a live example of these commands (slides 23-24).

# Manually Load and Configure a Sampler Plugin

## Additional Exercise Goals:

- Basic sampler plugin operation
  - Manual dynamic configuration using the “ldmsd\_controller” utility
  - Static configuration using a configuration file
  - man pages
    - man Plugin\_meminfo – opens meminfo plugin man pages
    - man Plugin\_vmstat – opens vmstat plugin man pages
- Use of ldms\_ls utility as a diagnostic tool
  - man pages
    - man ldms\_ls – opens ldms\_ls man pages

# LDMS Plugin Architecture



# Configure LDMS Daemon Sampler Plugin

## Goals:

- Load the “meminfo” sampler plugin
- Configure loaded “meminfo” sampler plugin
  - Give the set name (instance)
  - Give the node name (producer)
  - Give the component ID
  - Plugin-specific arguments
- Start sampler plugin with a particular sampling interval and offset

optional



# Connect Idmsd\_controller To An Idmsd

Set up “Idmsd\_controller” connection to the aggregator

```
$ldmsd_controller --host localhost --port 10001
```

```
welcome to the LDMSD control processor
```

```
sock:localhost:10001> help
```

See “LDMS HELP” slides starting at slide 77 for help results

- Note 1: The prompt tells you <transport>:<hostname>:<port>
- Note 2: You can use “quit” or Ctrl-d to **exit** or Ctrl-c to **kill** the Idmsd\_controller

\*A live example of these commands can be found here:

[LDMSD Controller Interface Video](#)

# Interactive Configuration Using The Idmsd\_controller

- Load the “meminfo” sampler plugin:

```
sock:localhost:10001> load name=meminfo
```

- Configure the “meminfo” sampler plugin:

```
sock:localhost:10001> config name=meminfo producer=<$HOSTNAME>  
instance=<$HOSTNAME>/meminfo component_id=<host number>
```

- EXAMPLE:

```
sock:localhost:10001> load name=meminfo *enter*  
sock:localhost:10001> config name=meminfo producer=ovis-demo-01  
instance=ovis-demo-01/meminfo component_id=1 *enter*
```

**producer:** Initialize the name of the sampler

**Instance:** Initialize the name of the node the sampler is running on

**component\_id:** Initialize with a number

# Query Current Sets On An LDMS Daemon Using “ldms\_ls”

- Use ldms\_ls to query the current sets available on an LDMS daemon

```
$ ldms_ls -h localhost -x sock -p 10001
```

ovis-demo-01/meminfo

# Get The Set Information Before Starting The “meminfo” Sampler Plugin

```
$ ldms_ls -h localhost -x sock -p 10001 -v ovis-demo-01/meminfo
```

Schema	Instance	Flags	Msize	Dsize	UID	GID	Perm	Update	Duration	Info
-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----
meminfo	ovis-demo-01/meminfo	L	1952	416	596	742	-rwxrwxrwx	0.000000	0.000000	
"updt_hint_us"="1000000:0"										
-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----

Total Sets: 1, Meta Data (kB): 1.95, Data (kB) 0.42, Memory (kB): 2.37

**NOTE:** The “[ovis-demo-01/meminfo](#)” is optional. It is suggested as it will be easier to identify certain sampler daemons when multiple are running on the same host and port.



## **EXAMPLE:** Interactive Configuration Using The ldmsd\_controller

Please see [Configuration Using LDMSD Controller Interface](#) to view a live example of these commands (slides 30-32).

# Query Current Metric Values Before Starting The “meminfo” Sampler Plugin



```
$ ldms_ls -x sock -p 10001 -l ovis-demo-01/meminfo
```

ovis-demo-01/meminfo: inconsistent, last update: Wed Dec 31 17:00:00 1969 -0700 [0us]

M u64	component_id	62
D u64	job_id	0
D u64	app_id	0
D u64	MemTotal	0
D u64	MemFree	0
D u64	MemAvailable	0
D u64	Buffers	0
D u64	Cached	0
D u64	SwapCached	0
D u64	Active	0
D u64	Inactive	0
D u64	Active(anon)	0
D u64	Inactive(anon)	0

- *Set is “inconsistent”*
- *Values have not yet been collected*

# Start The “meminfo” Sampler Plugin

- Start the “meminfo” sampler with a 1 second interval

```
sock:localhost:10001> start name=meminfo interval=1000000  
offset=0
```

- This starts the sampler updating the metric values every 1,000,000 micro-seconds = 1 second
- **Note 1:** “offset” defines micro-seconds after the second
- **Note 2:** If offset is not specified the timer starts when the sampler starts

# Query Current Metric Values After Starting The “meminfo” Sampler Plugin

```
$ ldms_ls -x sock -p 10001 -l ovis-demo-01/meminfo
```

ovis-demo-01/meminfo: consistent, last update: Tue Oct 08 17:52:45 2019 -0600 [2058us]

M u64	component_id	62
D u64	job_id	0
D u64	app_id	0
D u64	MemTotal	131899768
D u64	MemFree	129843340
D u64	MemAvailable	129364708
D u64	Buffers	20076
D u64	Cached	458024
D u64	SwapCached	0
D u64	Active	184380
D u64	Inactive	393140
D u64	Active(anon)	125324
D u64	Inactive(anon)	284684

- *Set is “consistent”*
- *Values have been collected*

# Periodically Re-Query Sampler and Run “memeater”

```
$ while true; do ldms_ls -h localhost -x sock -p 10001 -l | grep “Active ”; sleep 1; done
```

D u64	<b>Active</b>	192308
D u64	<b>Active</b>	192308
D u64	<b>Active</b>	191884
D u64	<b>Active</b>	192396
D u64	<b>Active</b>	192444
D u64	<b>Active</b>	192420
D u64	<b>Active</b>	192528
D u64	<b>Active</b>	192516

- *Note how the values change without/with “memeater” running*

In a separate terminal window, run the “memeater” executable to see both timestamps and values change:

```
$ /home/<user>/memeater/memeater
```

Note: You can edit and re-compile to change the allocation amounts and sleep time to adjust the rate of change.

# Check Source (/proc/meminfo) For Reference

```
$ cat /proc/meminfo
```

MemTotal:	131899768 kB
MemFree:	129828892 kB
MemAvailable:	129350280 kB
Buffers:	20076 kB
Cached:	458076 kB
SwapCached:	0 kB
Active:	192340 kB
Inactive:	393064 kB
Active(anon):	133212 kB
Inactive(anon):	284680 kB
Active(file):	59128 kB

## EXAMPLE: “meminfo” Sampler Plugin

Please see [Meminfo Sampler Daemon](#) to view a live example of these commands (slides 34-38).

# Dynamically Change The Sampling Interval

- Using `ldmsd_controller`, stop the plugin:

```
sock:localhost:10001> stop name=meminfo
```

Note: Querying with `ldms_ls` will show that the sampler is not updating

Note: We are still using the same sampler daemon from earlier. It should not be killed yet.

- Restart the plugin with a different (5 sec) interval:

```
sock:localhost:10001> start name=meminfo interval=5000000  
offset=0
```

Note: Querying with `ldms_ls` will show that the metric set is now updating only every five seconds

(More on dynamic configuration and resilience in Exercise 3)



# Kill Currently Running Daemons

- Kill all of your ldmsd in preparation for the next section

```
$ killall ldmsd
```

- Kill a particular ldmsd

```
$ ps auxw | grep ldmsd | grep -v grep
$ ovis_pu+ 3582 0.0 0.1 401604 2204 ? Ss1
12:51 0:00 ldmsd -x sock:10001 -S samplerd.sock
$ kill 3582
```

- Check to make sure it is dead

```
$ ps auxw | grep ldmsd | grep -v grep
```

## EXAMPLE: Change Sample Interval

Please see [Change Sample Interval for Meminfo](#) to view a live example of these commands (slides 40-41).

# Start a Idmsd and Sampler Plugin Using a Configuration File

- A Idmsd can be started using a configuration file
  - Syntax is identical to that used for manual configuration
  - Can be used to run and configure BOTH sampler and aggregator Idmsd
- Edit the sample configuration file, **as appropriate**, for the meminfo example:

```
$ cat /home/<user>/exercises/ldms/conf/simple_sampler.conf
```

- NOTE: If the “simple\_sampler.conf” is not there, then please create this file in this directory and populate it with the content below:

```
load name=meminfo
config name=meminfo producer=<$HOSTNAME> instance=<$HOSTNAME>/meminfo
component_id=<host number>
start name=meminfo interval=1000000
```

- Run an Idmsd using this configuration file (argument after the -c flag).  
Modify <user> to your user name.

```
$ ldmsd -x sock:10001 \  
-l /home/<user>/exercises/ldms/logs/sampler1.log \  
-c /home/<user>/exercises/ldms/conf/simple_sampler.conf
```

# Query The Metric Values: The “meminfo” Sampler Is Configured And Running

```
$ ldms_ls -x sock -p 10001 -l ovis-demo-01/meminfo
```

ovis-demo-01/meminfo: consistent, last update: Tue Oct 08 17:52:45 2019 -0600 [2058us]

M u64	component_id	62
D u64	job_id	0
D u64	app_id	0
D u64	MemTotal	131899768
D u64	MemFree	129843340
D u64	MemAvailable	129364708
D u64	Buffers	20076
D u64	Cached	458024
D u64	SwapCached	0
D u64	Active	184380
D u64	Inactive	393140
D u64	Active(anon)	125324
D u64	Inactive(anon)	284684

- *Set is “consistent”*
- *Values have been collected*

# Multiple Sampler Plugins

- Uncomment and **edit to reflect your host** the lines for the vmstat plugin in simple\_sampler.conf and restart the ldmsd.

```
load name=vmstat  
config name=vmstat producer=<hostname>  
instance=<hostname>/vmstat component_id=<hostnum>  
start name=vmstat interval=1000000 offset=0
```

- Query the ldmsd:

```
ldms_ls -h localhost -x sock -p 10001  
ovis-demo-01/vmstat  
ovis-demo-01/meminfo
```

Note: hostname is just a string and hostnum is just a uint\_64. Example:  
hostname=ovis-demo-01,  
hostnum=1

## EXAMPLE: Multiple Sampler Plugins

Please see [Multiple Plugin Sampler Daemon](#) to view a live example of these commands (slides 43-45).

# Configuration Tools Summary

## Dynamic/manual configuration (remote or local)

- `ldmsd_controller` – Python script that can connect to a `ldmsd` via a configured network socket or a local Unix Domain Socket (supports command completion)
- **`ldmsctl`** – C-based utility that can connect to a `ldmsd` via a configured network socket or a local Unix Domain Socket (doesn't support command completion)

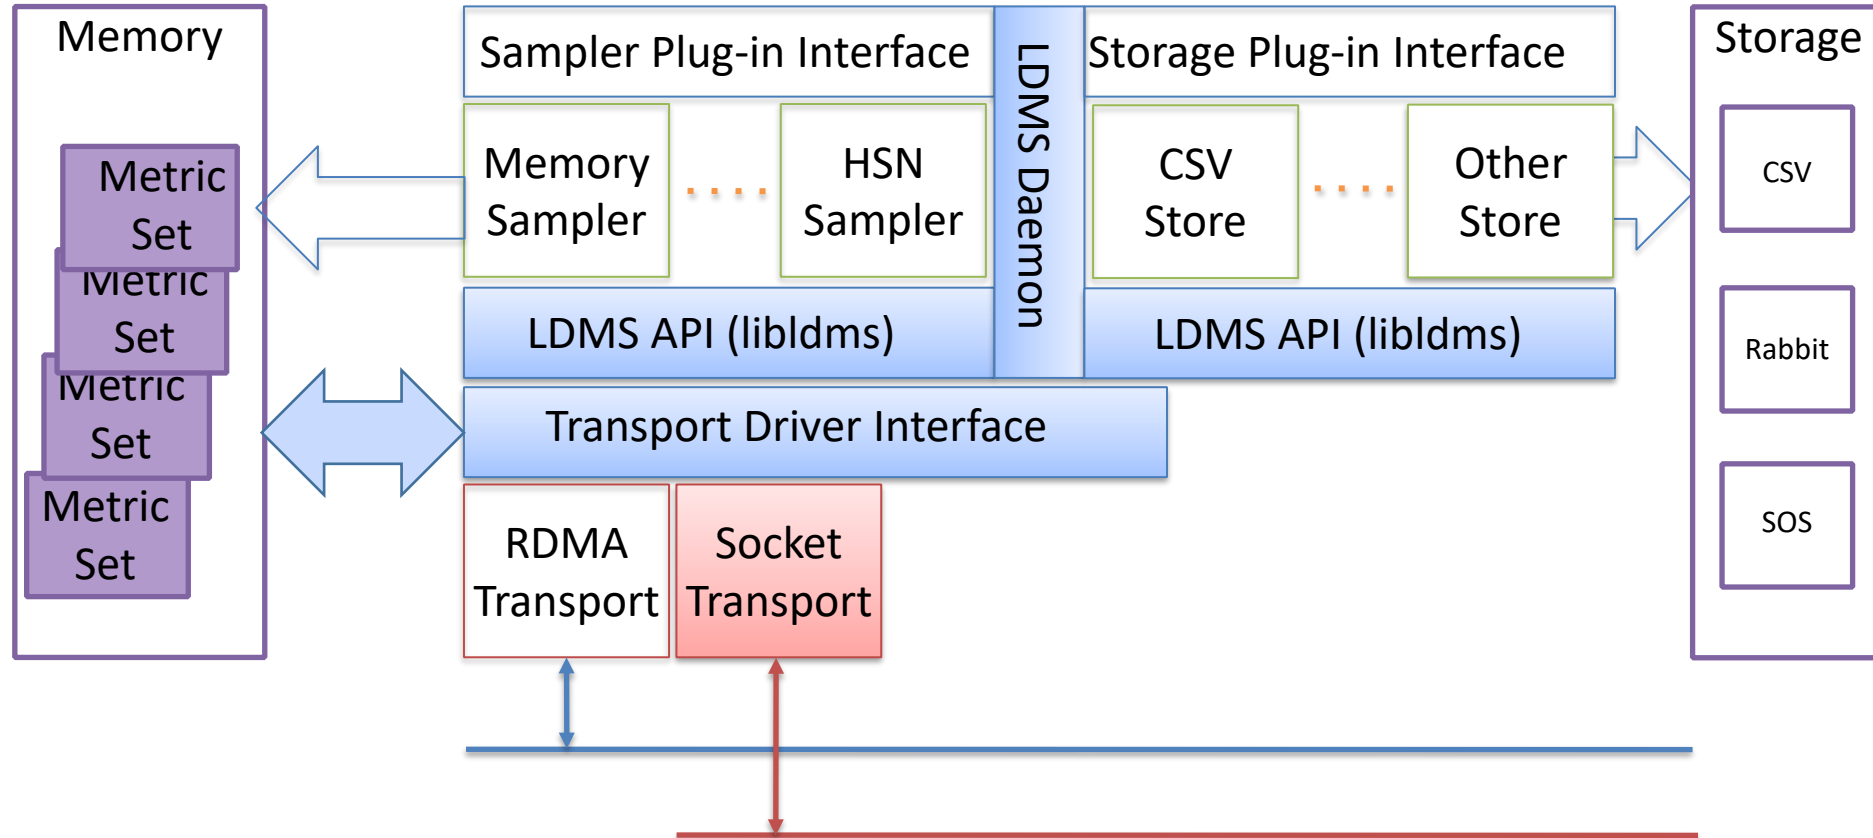
## Static configuration (local)

- Configuration file – loaded at `ldmsd` run time

# Exercise 3: Configure Aggregators



# LDMS Plugin Architecture



# Configure a LDMS daemon (ldmsd) to Aggregate metric set(s)



OGC

## Goals:

- Add list of connections to a ldmsd (connections to sampler ldmsd(s))
- Start the connections
- Create an “update policy”
  - Define an “update policy” update period
  - Define which sets an update policy refers to
- Start the “update policy”

# Start a Idmsd That Will Be Used For Aggregation

- (Re)start the sampler Idmsd from the previous exercise (can keep both meminfo and vmstat)
- Start new aggregator Idmsd with minimum configuration:

```
$ Idmsd -x sock:20001 -l /home/<user>/exercises/ldms/logs/agg1.log
```

- **-x:** Transport : listening port
- **-l:** Specify the log file path and name (this is not strictly necessary)

## NOTE:

- We will be using a different port number. Instead of 10001 we will be running a daemon on port 20001.
- Please refer to slides 41 & 42 for help in re-creating a sampler daemon

# Interactive Aggregator Configuration

- Set up “ldmsd\_controller” connection to the aggregator over socket

```
$ ldmsd_controller --host localhost --port 20001
```

```
Welcome to the LDMSD control processor
```

```
sock:localhost:20002>
```

# Simple Aggregator Producer Configuration

- Configure the aggregator to aggregate the “meminfo” set from your sampler daemon (listening on port 10001)

```
sock:localhost:20001> prdcr_add name=prdcr1 host=$HOSTNAME port=10001  
xprt=sock type=active interval=20000000  
sock:localhost:20001> prdcr_start name=prdcr1
```

- **name**: policy tag (this is just a string)
- **host**: hostname for the sampler daemon (e.g. ovis-demo-01)
- **port**: Listener port of the sampler daemon
- **xprt**: Transport the sampler daemon listens on
- **type**: Always “active” (the aggregator will initiate the connection with the sampler)
- **interval**: Re-connect interval (*not aggregation interval*)

# Check Aggregator Status

(**after producer** (prdcr) is started but **before** the **updater** (updtr) is started)

```
sock:localhost:20001> status
```

Name	Type	Interval	Offset	Libpath
------	------	----------	--------	---------

Name	Host	Port	Transport	State
------	------	------	-----------	-------

prdcr1	nid00062	10001	sock	CONNECTED
--------	----------	-------	------	-----------

nid00062/meminfo meminfo\_x86\_ven0000fam0006mod003F START

nid00062/vmstat vmstat\_x86\_ven0000fam0006mod003F START

Name	Interval	Offset	Mode	State
------	----------	--------	------	-------

Name	Container	Schema	Plugin	State
------	-----------	--------	--------	-------

# Query Current Metric Values On The Aggregator

```
$ ldms_ls -h localhost -x sock -p 20001 -l
```

```
$
```

**Note:** While status (previous slide) shows that the aggregator knows what sets the producer has, the `ldms_ls` query returns nothing because the updater had not yet been run and the set has not been populated with data.

## **EXAMPLE:** Simple Aggregator Producer Configuration

Please see the simple [Aggregator Producer Configuration](#) to view a live example of these commands (slides 51-55).



# Simple Aggregator Updater Configuration

- Configure the aggregator to **update** the “meminfo” set

```
sock:localhost:20001> updtr_add name=updtr1 interval=1000000  
offset=200000  
sock:localhost:20001> updtr_prdcr_add name=updtr1 regex=.*  
sock:localhost:20001> updtr_start name=updtr1
```

- **name:** policy tag (string)
- **interval:** update (pull) interval (in usec)
  - Example: interval=1000000 means pull data from sampler every 1 seconds
- **offset:** Target (in us) from <epoc sec>.000000
  - Example: offset=10000 means aggregate every <interval> seconds at 10ms into the second.
- **regex:** regular expression to match the target producers tag(s)

# Check Aggregator Status

(after starting both producer (prdcr) and updater (updtr) policies)

```
sock:localhost:20001> status
```

Name	Type	Interval	Offset	Libpath
------	------	----------	--------	---------

Name	Host	Port	Transport	State
------	------	------	-----------	-------

prdcr1	nid00062	10001	sock	CONNECTED
--------	----------	-------	------	-----------

nid00062/meminfo meminfo\_x86\_ven0000fam0006mod003F READY

nid00062/vmstat vmstat\_x86\_ven0000fam0006mod003F READY

Name	Interval	Offset	Mode	State
------	----------	--------	------	-------

updtr1	1000000	0	Pull	RUNNING
--------	---------	---	------	---------

prdcr1	nid00062	10001	sock	CONNECTED
--------	----------	-------	------	-----------

Name	Container	Schema	Plugin	State
------	-----------	--------	--------	-------

# Query Current Metric Values On The Aggregator

```
$ Idms_ls -h localhost -x sock -p 20001 -l ovis-demo-01/meminfo
```

nid00062/meminfo: consistent, last update: Wed Oct 09 18:30:49 2019 -0600 [2093us]

M u64	component_id	62
D u64	job_id	0
D u64	app_id	0
D u64	MemTotal	131899768
D u64	MemFree	129834752
D u64	MemAvailable	129356628
D u64	Buffers	20228
D u64	Cached	458892
D u64	SwapCached	0
D u64	Active	196708
D u64	Inactive	393768
D u64	Active(anon)	137336

# Check To See That Metrics Change In Both Samplers and Aggregators

- In a third window, run the mem eater executable to see changes in the dataset values in both the samplers and aggregators:

```
$ /home/<user>/memeater/memeater
```

- Check sampler using ldms\_ls:

```
$ while true; do ldms_ls -h localhost -x sock -p 10001 -l | grep "Active "; sleep 1; done
```

- Check aggregator using ldms\_ls:

```
$ while true; do ldms_ls -h localhost -x sock -p 20001 -l | grep "Active "; sleep 1; done
```

## **EXAMPLE:** Simple Aggregator Updater Configuration

Please see [Aggregator Updater Configuration](#) to view a live example of these commands (slides 57-60).

# Start Idmsd and Aggregation Using a Configuration File

- A Idmsd for performing aggregation can be started using a configuration file in the same manner as a Idmsd for sampling
- Configuration file syntax is identical to that used for manual configuration
- *Edit your* sample configuration file *to reflect your host*:

```
$ cat /home/<user>/exercises/ldms/conf/simple_agg.conf
```

**NOTE:** If the “simple\_agg.conf” is not there, then please create this file in this directory and populate it with the content below:

```
prdcr_add name=prdcr1 host=$HOSTNAME port=10001 xprt=sock type=active interval=20000000
prdcr_start name=prdcr1
updtr_add name=updtr1 interval=1000000 offset=200000
updtr_prdcr_add name=updtr1 regex=.*
updtr_start name=update_all
```

- Kill your aggregator and restart your aggregator Idmsd using this configuration file

```
$ ldmsd -x sock:20001 -l /home/<user>/exercises/ldms/log/aggd.log \
-c /home/<user>/exercises/ldms/conf/simple_agg.conf
```

# Query Current Metric Values On The Aggregator

```
$1dms_1s -x sock -p 20001 -l ovis-demo-01/meminfo
```

Ovis-demo-01/meminfo: consistent, last update: Wed Oct 09 18:30:49 2019 -0600 [2093us]

M u64	component_id	62
D u64	job_id	0
D u64	app_id	0
D u64	MemTotal	131899768
D u64	MemFree	129834752
D u64	MemAvailable	129356628
D u64	Buffers	20228
D u64	Cached	458892
D u64	SwapCached	0
D u64	Active	196708
D u64	Inactive	393768

## **EXAMPLE:** Simple Aggregator with Configuration File

Please see [Aggregator With Configuration File](#) to view a live example of these commands (slides 62-63).



# Exercise 4: Aggregating From Remote Hosts: Building a Distributed Monitoring System

# Aggregate From All Student VMs

- Kill aggregator ldmsd
- Edit /home/<user>/exercises/ldms/conf/agg.conf:

```
$ cat /home/<user>/exercises/ldms/conf/agg.conf
```

```
#prdcr_add name=prdcr1 type=active host=compute1 port=10001 xprt=sock interval=20000000
```

```
#prdcr_add name=prdcr2 type=active host=compute2 port=10001 xprt=sock interval=20000000
```

```
#prdcr_add name=prdcr3 type=active host=compute3 port=10001 xprt=sock interval=20000000
```

```
prdcr_start_regex regex=.*
```

*START (connect to) ALL PRODUCERS*

```
updtr_add name=updtr1 interval=1000000 offset=200000
```

*UPDATE AT 1 SEC INTERVALS*

```
updtr_prdcr_add name=updtr1 regex=.*
```

DO THIS ON ALL PRODUCERS

```
updtr_match_add name=updtr1 match=schema regex=meminfo
```

*RESTRICT TO SETS WITH schema=meminfo*

```
updtr_start name=updtr1
```

START UPDATER POLICY "updtr1"

```
updtr_add name=updtr2 interval=2000000 offset=200000
```

*UPDATE AT 2 SECOND INTERVALS*

```
updtr_prdcr_add name=updtr2 regex=.*
```

DO THIS ON ALL PRODUCERS

```
updtr_match_add name=updtr2 match=schema regex=vmstat
```

*RESTRICT TO SETS WITH schema=vmstat*

```
updtr_start name=updtr2
```

START UPDATER POLICY "updtr1"

# Aggregate From All Student VMs (cont'd)

- Restart Idmsd using your edited configuration file

```
$ Idmsd -x sock:20001 -l /home/<user>/exercises/ldms/log/aggd.log \  
-c /home/<user>/exercises/ldms/conf/agg.conf
```

**NOTE:** If the “**agg.conf**” has not yet been populated, then create this file by first copying the “simple\_aggregator.conf” and adding the content from the previous slide at the end of the file.

LDMS supports complex topologies:

- Multiple Idmsd (aggregators) can pull from the same Idmsd (sampler or aggregator)
- Can daisy chain aggregators
  - Hierarchical
  - Support both fan-in and fan-out topologies

# Check Aggregator Status

```
sock:localhost:20001> status
```

Name	Type	Interval	Offset	Libpath
------	------	----------	--------	---------

Name	Host	Port	Transport	State
------	------	------	-----------	-------

prdcr1	nid00062	10001	sock	CONNECTED
nid00062/meminfo meminfo_x86_ven0000fam0006mod003F READY				
nid00062/vmstat vmstat_x86_ven0000fam0006mod003F READY				
prdcr2	nid00063	10001	sock	CONNECTED
nid00063/meminfo meminfo_x86_ven0000fam0006mod003F READY				
nid00063/vmstat vmstat_x86_ven0000fam0006mod003F READY				

prdcr3	nid00012	10001	sock	DISCONNECTED
--------	----------	-------	------	--------------

Name	Interval	Offset	Mode	State
------	----------	--------	------	-------

updtr1	1000000	200000	Pull	RUNNING
--------	---------	--------	------	---------

prdcrl	nid00062	10001	sock	CONNECTED
--------	----------	-------	------	-----------

prdcrl	nid00063	10001	sock	CONNECTED
--------	----------	-------	------	-----------

prdcrl	nid00012	10001	sock	DISCONNECTED
--------	----------	-------	------	--------------

Name	Container	Schema	Plugin	State
------	-----------	--------	--------	-------

# Using The Distributed System

- **Exercise** - Loop `ldms_ls` while running your `memeater` executable – see your and others data values change

```
$ while true; do ldms_ls -h localhost -x sock -p 20001 -l | grep "Active " -B9; sleep 1; done
```

Explore basic dynamic configuration changes and resilience in the next exercise

# EXAMPLE: Aggregate from Multiple VMs

Please see [Aggregate From Multiple VMs](#) to view a live example of these commands (slides 66-69) .

# Exercise 5: Basic Dynamic Configurations and Resilience

# Basic Dynamic Configuration Changes

- Dynamic configuration
  - Sampler daemons (from exercise 1 slide 40)
    - Stopping sampler plugins
    - Starting sampler plugins with different intervals
  - Aggregator daemons
    - Automatic detection of new metric sets on connected sampler ldmsd
    - Stopping producer (prdcr) and updater (updtr) policies
    - Changing updater intervals



# Dynamically Changing a Sampler Plugin's Interval Parameters (also exercise 1 slide 40)

- Using `ldmsd_controller`, stop the plugin:

```
sock:localhost:10001> stop name=meminfo
```

Note: Querying with `ldms_ls` will show that the sampler has stopped

- Restart the plugin with a different interval:

```
sock:localhost:10001> start name=meminfo  
interval=5000000 offset=0
```

Note: Querying with `ldms_ls` will show that the metric set is now updating only every five seconds

# Dynamic Changes and Aggregator Robustness OGC

- On-the-fly additions of samplers will be discovered by the aggregating Idmsd
  - **Exercise** – one student will add the vmstat sampler, using Idmsd\_controller, to their running sampler Idmsd.
    - All others will see it, using Idms\_ls, appear in their aggregators which are pulling from that sampler.
  - **Exercise** – one student will stop their meminfo sampler, using Idmsd\_controller, on their running sampler Idmsd.
    - All others will see, using Idms\_ls, that the timestamp in that student's metric set ceases to update.
  - **Exercise** – the same student will restart their meminfo sampler, using Idmsd\_controller, on their running Idmsd.
    - All others will see, using Idms\_ls, that the timestamp in that student's metric set resumes updating.

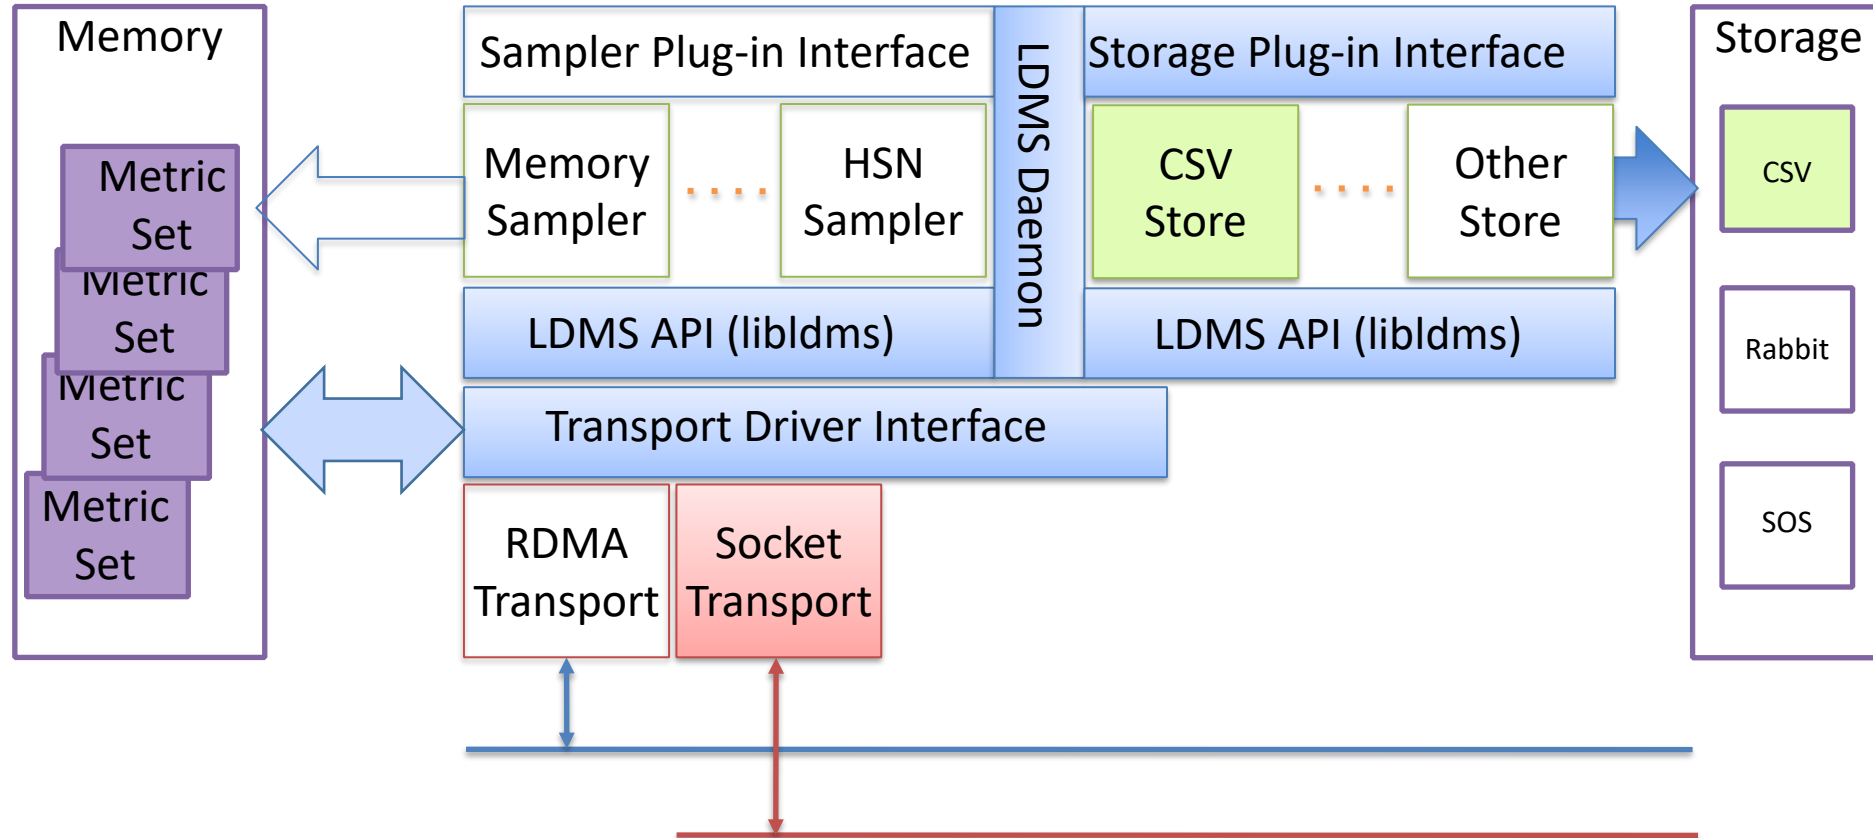
# Dynamic Changes and Robustness (cont'd)



- Samplers and Aggregators can be started in any order
  - **Exercise** – Use your modified configuration files to start the aggregator ldmsd before starting the sampler ldmsd
    - Use ldms\_ls to convince yourself that, whether a sampler ldmsd is started before or after an aggregator ldmsd, you are able to see the data generated at the sampler ldmsd on the aggregator ldmsd when both are running
- LDMS collection and transport are robust to Samplers and Aggregators being killed and restarted
  - **Exercise** – one student will kill their sampler ldmsd. All other students will see from ldms\_ls timestamp that the student's metric set is removed from the list.
  - **Exercise** – the same student will restart their sampler ldmsd. All other students will see from ldms\_ls timestamp that the metric set reappears and resumes updating (after up to the producer reconnect interval of 20 seconds).
  - **Exercise** – Each student will stop and re-start their aggregator ldmsd and verify, using ldms\_ls, that they are able to see appropriate data.

# Exercise 6: Storing Data In CSV Format

# LDMS Plugin Architecture



# Storing Data: CSV Store Plugin

- Goals:
  - Configure an aggregator Idmsd with a CSV store plugin using Idmsd\_controller
  - Configure an aggregator Idmsd with a CSV store plugin using a configuration file
  - Minimal store options (don't buffer data)
- Example output from the “meminfo” sampler:

#Time,Time\_usec,ProducerName,component\_id,job\_id,MemTotal,MemFree,MemAvailable,Buffers,Cached,SwapCached,Active,Inactive,Active(anon),Inactive(anon),Active(file),Inactive(file),Unevictable,Mlocked,SwapTotal,SwapFree,Dirty,Writeback,AnonPages,Mapped,Shmem,Slab,SReclaimable,SUnreclaim,KernelStack,PageTables,NFS\_Unstable,Bounce,WritebackTmp,CommitLimit,Committed\_AS,VmallocTotal,VmallocUsed,VmallocChunk,HardwareCorrupted,AnonHugePages,HugePages\_Total,HugePages\_Free,HugePages\_Rsvd,HugePages\_Surp,Hugepagesize,DirectMap4k,DirectMap2M

1487105964.002482,2482,ovis-demo-09,9,  
0,1884188,571028,1688632,0,1212004,6108,104536,1122496,8276,8580,96260,1113916,0,0,839676,793956,420,0,10552,24812,17  
96,52124,40104,12020,1792,3280,0,0,0,1781768,387984,34359738367,7216,34359728128,0,2048,0,0,0,0,2048,47040,2050048

1487105963.002583,2583,ovis-demo-02,2,  
0,1884188,1665280,1671132,948,107512,0,71540,80920,44128,8308,27412,72612,0,0,839676,839676,0,0,44000,22264,8436,35680  
,24304,11376,1600,2940,0,0,0,1781768,296444,34359738367,7216,34359728128,0,6144,0,0,0,0,2048,34752,2062336

1487105963.001964,1964,ovis-demo-08,8,  
0,1884188,1623168,1644996,948,129700,0,89312,101956,60788,8332,28524,93624,0,0,839676,839676,0,0,60620,23912,8500,3645  
6,24608,11848,1872,4364,0,0,0,1781768,403252,34359738367,7216,34359728128,0,16384,0,0,0,0,2048,44992,2052096

# CSV Store: Manual Aggregator Configuration

- Configure the aggregator to **store** the “meminfo” set to a **CSV** file using `ldmsd_controller`
  - Create a directory for the CSV data
  - Load the `store_csv` plugin
  - Configure the plugin

```
$ mkdir -p /home/<user>/exercises/ldms/data/CSV
$ ldmsd_controller --host localhost --port 20001
sock:localhost:20001> load name=store_csv
sock:localhost:20001> config name=store_csv path=/home/<user>/exercises/ldms/data/CSV buffer=0
```

- Check status

```
sock:localhost:20001> status
Name      Type      Interval  Offset  Libpath
-----
csv       store      1000000   0       /home/<user>/Build/OVIS-4.3.1/lib/ovis-ldms/libstore_csv.so
```

- **name**: plugin name
- **path**: Path to the base directory for the csv file container. This directory must pre-exist.
- **buffer**: ‘0’ to disable buffering **# USE WITH CAUTION!**
- **man page**:
  - `man Plugin_store_csv` – opens `store_csv` plugin man pages

# CSV Store: Manual Aggregator Configuration (cont.)



- Configure the aggregator to **store** the “meminfo” set to a csv file.

```
sock:localhost:20001> strgp_add name=meminfo_store_csv plugin=store_csv  
container=memory_metrics schema=meminfo
```

- Check status

```
sock:localhost:20001> status
```

Name	Container	Schema	Plugin	State
-----	-----	-----	-----	-----
meminfo_store_csv	memory_metrics	meminfo	store_csv	STOPPED

producers:  
metrics:

- **name**: storage policy tag
- **plugin**: store plugin used for storing metric set data
- **container**: the storage backend container name. For csv, this is the directory where the output file will go. This will be created.
- **schema**: metric set schema to be stored



# CSV Store: Manual Aggregator Configuration (cont.)



```
sock:localhost:20001> strgp_start name=meminfo_store_csv
```

- Check status

```
sock:localhost:20001> status
```

Name	Container	Schema	Plugin	State
meminfo_store_csv	memory_metrics	meminfo	store_csv	RUNNING

producers:

metrics: component\_id job\_id app\_id MemTotal MemFree MemAvailable Buffers Cached  
SwapCached Active Inactive Active(anon) Inactive(anon) Active(file) Inactive(file) Unevictable  
Mlocked SwapTotal SwapFree Dirty Writeback AnonPages Mapped Shmem Slab SReclaimable  
SUnreclaim KernelStack PageTables NFS\_Unstable Bounce WritebackTmp CommitLimit  
Committed\_AS VmallocTotal VmallocUsed VmallocChunk HardwareCorrupted HugePages\_Total  
HugePages\_Free HugePages\_Rsvd HugePages\_Surp Hugepagesize DirectMap4k DirectMap2M  
DirectMap1G

- **name:** storage policy tag

# CSV Store: LDMSD Status

```
sock:localhost:20001> status
```

Name	Type	Interval	Offset	Libpath
csv	store	1000000	0	/home/<user>/Build/OVIS-4.3.1/lib/ovis-ldms/libstore_csv.so

Name	Host	Port	Transport	State
------	------	------	-----------	-------

prdc1	nid00052	10001	sock	CONNECTED
	nid00052/meminfo		meminfo	READY
	nid00052/vmstat		vmstat	READY
prdc2	nid00053	10001	sock	CONNECTED
	nid00053/meminfo		meminfo	READY
	nid00053/vmstat		vmstat	READY

Name	Interval	Offset	Mode	State
updtr1	1000000	200000	Pull	RUNNING
prdc1	nid00052	10001	sock	CONNECTED
prdc2	nid00053	10001	sock	CONNECTED

Name	Container	Schema	Plugin	State
meminfo_store_csv	memory_metrics	meminfo	store_csv	RUNNING

producers:

metrics: component id job\_id app\_id MemTotal MemFree MemAvailable Buffers Cached SwapCached Active Inactive Active(anon)  
Inactive(anon) Active(file) Inactive(file) Unevictable Mlocked SwapTotal SwapFree Dirty Writeback AnonPages Mapped Shmem Slab  
SReclaimable SUnreclaim KernelStack PageTables NFS Unstable Bounce WritebackTmp CommitLimit Committed\_AS VmallocTotal VmallocUsed  
VmallocChunk HardwareCorrupted HugePages\_Total HugePages\_Free HugePages\_Rsvd HugePages\_Surp Hugepagesize DirectMap4k  
DirectMap2M DirectMap1G

# Examining The CSV File

## Exercise: Check the CSV file

```
$ head /home/<user>/exercises/ldms/data/CSV/memory_metrics/meminfo  
$ tail -f /home/<user>/exercises/ldms/data/CSV/memory_metrics/meminfo
```

- If aggregating from others' vm's, you will see multiple hosts in the output

## Exercise: View data changes:

- Run the mem eater executable

```
$ /home/<user>/exercises/ldms/code/memeater
```

- Compare the live memeater output using tail -f ".../meminfo" values

## EXAMPLE: CSV Store - Manual Aggregator Configuration

Please see [Manual CSV Store](#) to view a live example of these commands (slides 78-83).

# CSV Store: Start and Configure Aggregator Using a Configuration File

- Edit aggregator configuration file, as appropriate, at:  
`/home/<user>/exercises/ldms/conf/agg.conf`

```
load name=store_csv
config name=store_csv path=/home/<user>/exercises/ldms/data/CSV buffer=0
strgp_add name=meminfo-store_csv plugin=store_csv container=memory_metrics
schema=meminfo
strgp_start name=meminfo-store_csv
```

- Restart your aggregator using:  
`/home/<user>/exercises/ldms/scripts/start_agg.conf`

## **EXAMPLE:** CSV Store - Start and Configure Aggregator Using a Configuration File

Please see [CSV Store Using Configuration File](#) to view a live example of these commands (slide 85).

Basics End

# LDMS HELP



# Idms\_controller: “help” Topics

```
sock:localhost.10001> help
```

Documented commands (type help <topic>):

=====

EOF	logrotate	setgroup_del	term
comment	oneshot	setgroup_ins	udata
<b>config</b>	plugn_sets	setgroup_mod	udata_regex
connect	<b>prdcr_add</b>	setgroup_rm	<b>updtr_add</b>
daemon_exit	prdcr_del	shell	updtr_del
daemon_status	prdcr_hint_tree	source	updtr_match_add
env	prdcr_set_status	<b>start</b>	updtr_match_del
failover_config	<b>prdcr_start</b>	stop	updtr_prdcr_add
failover_peercfg_start	prdcr_start_regex	<b>strgp_add</b>	updtr_prdcr_del
failover_peercfg_stop	prdcr_status	strgp_del	<b>updtr_start</b>
failover_start	prdcr_stop	strgp_metric_add	updtr_status
failover_status	prdcr_stop_regex	strgp_metric_del	updtr_stop
failover_stop	prdcr_subscribe	strgp_prdcr_add	updtr_task
greeting	publish	strgp_prdcr_del	<b>usage</b>
<b>help</b>	quit	<b>strgp_start</b>	version
include	say	strgp_status	
<b>load</b>	set_route	strgp_stop	
loglevel	setgroup_add	subscribe	

Undocumented commands:

=====

example	plugn_status	script	status	try
---------	--------------	--------	--------	-----

Definitely use for samplerd  
Definitely use for aggregators  
Definitely use for aggregators that store  
Use to load and config plugin  
Get help and daemon status

# ldmsd\_controller: Command Help

```
sock:localhost:10001> help prdcr_add
```

Add an LDMS Producer to the Aggregator

Parameters:

name=	A unique name for this Producer
xprt=	The transport name [sock, rdma, ugni]
host=	The hostname of the host
port=	The port number on which the LDMS is listening
type=	The connection type [active, passive]
interval=	The connection retry interval (us)